

# CSCI 4710/5710 Databases

## SQL JOINS





# Why JOINS

- Joins are used to select data from multiple tables.
  - Display the customer name with orders instead of the customer ID.
    - Can't display this information using one table.
  - Will need the order information.
    - From the orders table.
  - Will need the customer information.
    - From the customers table.



# JOINS – How to.

- Joins can be achieved in two different ways:

- Traditional method:

- Using conditions in the where clause.

- Example: Display ISBN, title, retail category and publisher's name for all books.

```
SELECT B.ISBN, B.Title, B.Retail, B.Category , P.Name as 'Published By'  
FROM BOOKS B , PUBLISHER P WHERE B.PubID = P.PubID;
```

| ISBN       | Title                          | Retail | Category    | Published By           |
|------------|--------------------------------|--------|-------------|------------------------|
| 0401140733 | REVENGE OF MICKEY              | 22.00  | FAMILY LIFE | PRINTING IS US         |
| 9247381001 | HOW TO MANAGE THE MANAGER      | 31.95  | BUSINESS    | PRINTING IS US         |
| 4981341710 | BUILDING A CAR WITH TOOTHPICKS | 59.95  | CHILDREN    | PUBLISH OUR WAY        |
| 9959789321 | E-BUSINESS THE EASY WAY        | 54.50  | COMPUTER    | PUBLISH OUR WAY        |
| 1915762492 | HANDCRANKED COMPUTERS          | 25.00  | COMPUTER    | AMERICAN PUBLISHING    |
| 3957136468 | HOLY GRAIL OF ORACLE           | 75.95  | COMPUTER    | AMERICAN PUBLISHING    |
| 8843172113 | DATABASE IMPLEMENTATION        | 55.95  | COMPUTER    | AMERICAN PUBLISHING    |
| 0132149871 | HOW TO GET FASTER PIZZA        | 29.95  | SELF HELP   | READING MATERIALS INC. |
| 0299282519 | THE WOK WAY TO COOK            | 28.75  | COOKING     | READING MATERIALS INC. |
| 1059831198 | BODYBUILD IN 10 MINUTES A DAY  | 30.95  | FITNESS     | READING MATERIALS INC. |
| 3437212490 | COOKING WITH MUSHROOMS         | 19.95  | COOKING     | READING MATERIALS INC. |
| 2147428890 | SHORTEST POEMS                 | 39.95  | LITERATURE  | REED-N-RITE            |
| 2491748320 | PAINLESS CHILD-REARING         | 89.95  | FAMILY LIFE | REED-N-RITE            |
| 8117949391 | BIG BEAR AND LITTLE DOVE       | 8.95   | CHILDREN    | REED-N-RITE            |



# JOINS – How to.

- Join Method:

- Using the JOIN keyword in the FROM clause.

```
SELECT B.ISBN, B.Title, B.Retail, B.Category , P.Name as 'Published By'  
FROM BOOKS B JOIN PUBLISHER P on B.PubID = P.PubID
```

| ISBN       | Title                          | Retail | Category    | Published By           |
|------------|--------------------------------|--------|-------------|------------------------|
| 0401140733 | REVENGE OF MICKEY              | 22.00  | FAMILY LIFE | PRINTING IS US         |
| 9247381001 | HOW TO MANAGE THE MANAGER      | 31.95  | BUSINESS    | PRINTING IS US         |
| 4981341710 | BUILDING A CAR WITH TOOTHPICKS | 59.95  | CHILDREN    | PUBLISH OUR WAY        |
| 9959789321 | E-BUSINESS THE EASY WAY        | 54.50  | COMPUTER    | PUBLISH OUR WAY        |
| 1915762492 | HANDCRANKED COMPUTERS          | 25.00  | COMPUTER    | AMERICAN PUBLISHING    |
| 3957136468 | HOLY GRAIL OF ORACLE           | 75.95  | COMPUTER    | AMERICAN PUBLISHING    |
| 8843172113 | DATABASE IMPLEMENTATION        | 55.95  | COMPUTER    | AMERICAN PUBLISHING    |
| 0132149871 | HOW TO GET FASTER PIZZA        | 29.95  | SELF HELP   | READING MATERIALS INC. |
| 0299282519 | THE WOK WAY TO COOK            | 28.75  | COOKING     | READING MATERIALS INC. |
| 1059831198 | BODYBUILD IN 10 MINUTES A DAY  | 30.95  | FITNESS     | READING MATERIALS INC. |
| 3437212490 | COOKING WITH MUSHROOMS         | 19.95  | COOKING     | READING MATERIALS INC. |
| 2147428890 | SHORTEST POEMS                 | 39.95  | LITERATURE  | REED-N-RITE            |
| 2491748320 | PAINLESS CHILD-REARING         | 89.95  | FAMILY LIFE | REED-N-RITE            |
| 8117949391 | BIG BEAR AND LITTLE DOVE       | 8.95   | CHILDREN    | REED-N-RITE            |



# Types of JOINS.

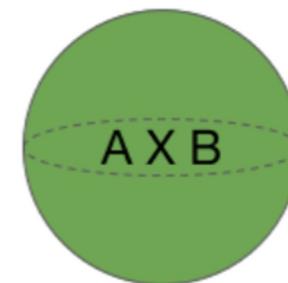
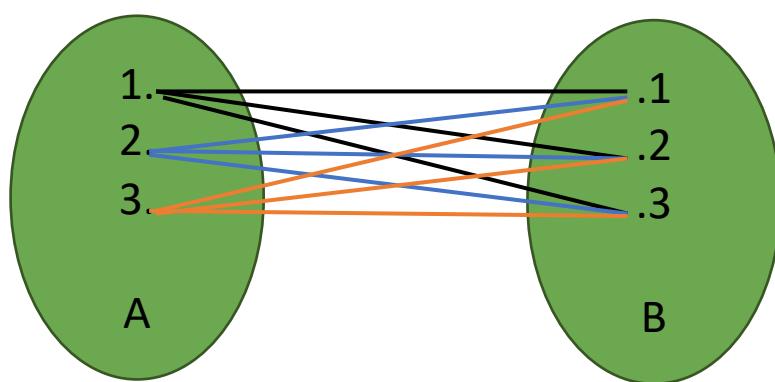
- Cartesian Join.
  - Also called cross join.
- Inner Join.
  - Equality join.
  - Non-equality join.
  - Self join.
- Outer Join.
  - Right Join.
  - Left Join.
  - Full Join.



SAINT LOUIS  
UNIVERSITY

# Cartesian (CROSS JOIN).

- Used to create all possible combinations between tables.
  - For two tables A and B, then :



CARTESIAN  
(CROSS) JOIN



SAINT LOUIS  
UNIVERSITY

# Cartesian (CROSS JOIN).

## Cartesian-Product Operation – Example

- Relations  $r, s$ :

| $A$      | $B$ |
|----------|-----|
| $\alpha$ | 1   |
| $\beta$  | 2   |

$r$

| $C$      | $D$ | $E$ |
|----------|-----|-----|
| $\alpha$ | 10  | +   |
| $\beta$  | 10  | +   |
| $\beta$  | 20  | -   |
| $\gamma$ | 10  | -   |

$s$

- $r \times s$

| $A$      | $B$ | $C$      | $D$ | $E$ |
|----------|-----|----------|-----|-----|
| $\alpha$ | 1   | $\alpha$ | 10  | +   |
| $\alpha$ | 1   | $\beta$  | 10  | +   |
| $\alpha$ | 1   | $\beta$  | 20  | -   |
| $\alpha$ | 1   | $\gamma$ | 10  | -   |
| $\beta$  | 2   | $\alpha$ | 10  | +   |
| $\beta$  | 2   | $\beta$  | 10  | +   |
| $\beta$  | 2   | $\beta$  | 20  | -   |
| $\beta$  | 2   | $\gamma$ | 10  | -   |

# Cartesian (CROSS JOIN).

- Let us try with an example by consider the following tables:

CARS

| Model         |
|---------------|
| BMW           |
| Ferrari       |
| Mercedes Benz |

COLORS

| ColorCode | ColorName |
|-----------|-----------|
| 1         | Red       |
| 2         | Blue      |
| 3         | Yellow    |
| 4         | Green     |

- Display all cars in all possible colors.



# Cartesian (CROSS JOIN).

- Can be done in two ways:

1. Omitting the WHERE clause:

```
SELECT * FROM CARS, COLORS order by Model, ColorCode;
```

2. Using CROSS JOIN:

```
SELECT * FROM CARS CROSS JOIN COLORS order by Model, ColorCode;
```

- Both will give the same result

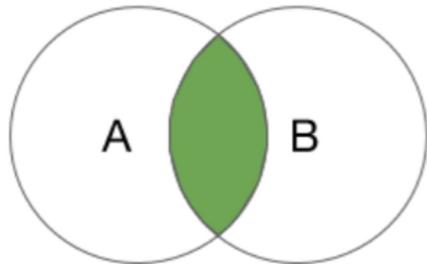
| Model         | ColorCode | ColorName |
|---------------|-----------|-----------|
| BMW           | 1         | Red       |
| BMW           | 2         | Blue      |
| BMW           | 3         | Yellow    |
| BMW           | 4         | Green     |
| Ferrari       | 1         | Red       |
| Ferrari       | 2         | Blue      |
| Ferrari       | 3         | Yellow    |
| Ferrari       | 4         | Green     |
| Mercedes Benz | 1         | Red       |
| Mercedes Benz | 2         | Blue      |
| Mercedes Benz | 3         | Yellow    |
| Mercedes Benz | 4         | Green     |



SAINT LOUIS  
UNIVERSITY

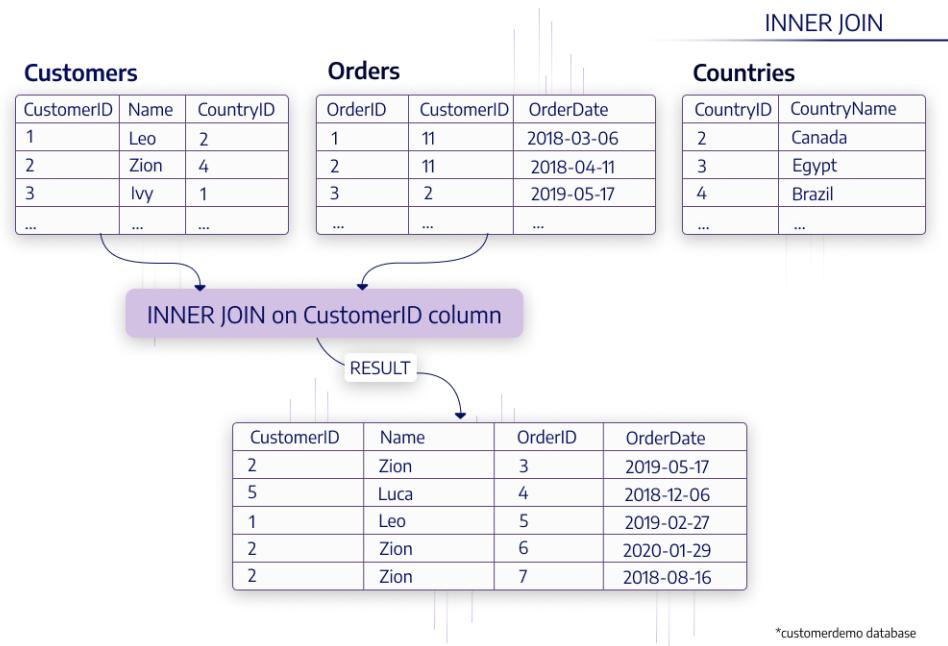
# Inner Join

- Basic idea:



## INNER JOIN

- Selects records that have matching values in both tables.
- Example: **SELECT Orders.OrderID, Customers.CustomerName  
FROM Orders  
INNER JOIN Customers ON Orders.CustomerID = Customers.CustomerID;**





SAINT LOUIS  
UNIVERSITY

# Inner Join

- The resulting table will contain all the attributes from both the tables including common column also.

**Example:** SELECT \* FROM student S INNER JOIN Marks M  
ON S.Roll\_No = M.Roll\_No;

| Roll_No | Name | Roll_No | Marks |
|---------|------|---------|-------|
| 2       | B    | 2       | 70    |
| 3       | C    | 3       | 50    |



SAINT LOUIS  
UNIVERSITY

# Equality Joins

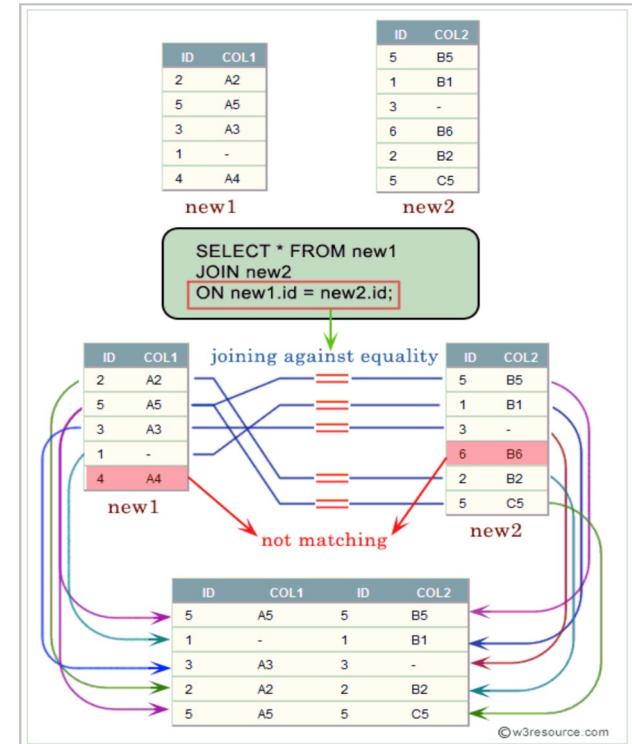
- Link rows through equivalent data that exists in both tables.

- Matching tables on joining condition

- Based on a matching key (PK-FK)

- Output: Intersection of two tables

- Only include those matched
  - Non-matched rows excluded





# Equality Joins

- Examples:
  - Display order numbers, dates and customer's full name for all orders.

```
SELECT OrderID, OrderDate , concat(FirstName, ' ', LastName) as 'Customer Name'  
FROM ORDERS, CUSTOMERS where ORDERS.CustomerID = CUSTOMERS.CustomerID
```

- Display order numbers, dates and customer's full name for all customers living in Florida.
- ```
SELECT OrderID, OrderDate , concat(FirstName, ' ', LastName) as 'Customer Name'  
FROM ORDERS, CUSTOMERS where ORDERS.CustomerID = CUSTOMERS.CustomerID  
AND CUSTOMERS.State = 'FL'
```

If there is same column name in more than one table???



# Column Ambiguity

- When the same column name is present in more than one table.
  - Considered ambiguous.
  - Which column does the query refer to?
- Resolved by specifying the table containing the column.
  - Example: the column “CustomerID” exists in both the ORDERS table and the CUSTOMERS Table:
  - Trying this query:

```
SELECT OrderID, OrderDate, CustomerID , concat(FirstName, ' ', LastName) as 'Customer Name'  
FROM ORDERS, CUSTOMERS where ORDERS.CustomerID = CUSTOMERS.CustomerID
```

Will result in an error: Error Code: 1052. Column 'CustomerID' in field list is ambiguous



# Column Ambiguity

- To resolve the issue:
  - Specify which table to pull data from:

```
SELECT OrderID, OrderDate, CUSTOMERS.CustomerID , concat(FirstName, ' ', LastName) as 'Customer Name'  
FROM ORDERS, CUSTOMERS where ORDERS.CustomerID = CUSTOMERS.CustomerID
```

- Or similarly:

```
SELECT OrderID, OrderDate, ORDERS.CustomerID , concat(FirstName, ' ', LastName) as 'Customer Name'  
FROM ORDERS, CUSTOMERS where ORDERS.CustomerID = CUSTOMERS.CustomerID
```



# Aliasing

- Joining more than two tables requires more join conditions.
- Column aliasing makes it easier to specify tables and remove ambiguity.
  - It can be done using an alias (usually a letter or two) to refer to a table.
  - Syntax:
    - In the from clause:
    - Follow the table name with the alias.
    - Use the alias in the select clause as:
      - Alias.ColumnName.



# Joining Multiple Tables

- Example: Display the names of the customers along with the titles of the books they have ordered.
  - This requires fetching data from multiple tables.



SAINT LOUIS  
UNIVERSITY

# Joining Multiple Tables

- Visualization:

CUSTOMERS

| Field      | Type        | Null | Key |
|------------|-------------|------|-----|
| CustomerID | int(11)     | NO   | PRI |
| LastName   | varchar(10) | NO   |     |
| FirstName  | varchar(10) | NO   |     |
| Address    | varchar(20) | YES  |     |
| City       | varchar(12) | YES  |     |
| State      | varchar(2)  | YES  |     |
| Zip        | varchar(5)  | YES  |     |
| Referred   | int(11)     | YES  |     |
| Region     | char(2)     | YES  |     |

ORDERS

| Field      | Type        | Null | Key |
|------------|-------------|------|-----|
| OrderID    | int(11)     | NO   | PRI |
| CustomerID | int(11)     | YES  | MUL |
| OrderDate  | date        | NO   |     |
| ShipDate   | date        | YES  |     |
| ShipStreet | varchar(18) | YES  |     |
| ShipCity   | varchar(15) | YES  |     |
| ShipState  | varchar(2)  | YES  |     |
| ShipZip    | varchar(5)  | YES  |     |
| ShipCost   | double(4,2) | YES  |     |

BOOKS

| Field    | Type        | Null | Key |
|----------|-------------|------|-----|
| ISBN     | varchar(10) | NO   | PRI |
| Title    | varchar(30) | YES  |     |
| PubDate  | date        | YES  |     |
| PubID    | int(11)     | YES  | MUL |
| Cost     | double(5,2) | YES  |     |
| Retail   | double(5,2) | YES  |     |
| Discount | double(4,2) | YES  |     |
| Category | varchar(12) | YES  |     |

ORDERITEMS

| Field    | Type        | Null | Key |
|----------|-------------|------|-----|
| OrderID  | int(11)     | NO   | PRI |
| ItemID   | int(11)     | NO   | PRI |
| ISBN     | varchar(10) | YES  | MUL |
| Quantity | int(11)     | NO   |     |
| PaidEach | double(5,2) | NO   |     |



# Joining Multiple Tables

- Example: Display the names of the customers along with the titles of the books they have ordered.
  - This requires fetching data from multiple tables.
    - Customer Name: Customers Table.
    - Book Title: Books Table:
    - Which books and in which order: Order Items Table.
    - Which order belongs to which customer: Orders table:

```
SELECT C.FirstName, C.LastName, B.Title FROM  
CUSTOMERS C, ORDERS O, ORDERITEMS OI, BOOKS B  
WHERE C.CustomerID = O.CustomerID  
AND O.OrderID = OI.OrderID  
AND OI.ISBN = B.ISBN;
```



SAINT LOUIS  
UNIVERSITY

# Natural Join

## Natural Join Operation – Example

- Relations  $r, s$ :

| $A$      | $B$ | $C$      | $D$ |
|----------|-----|----------|-----|
| $\alpha$ | 1   | $\alpha$ | a   |
| $\beta$  | 2   | $\gamma$ | a   |
| $\gamma$ | 4   | $\beta$  | b   |
| $\alpha$ | 1   | $\gamma$ | a   |
| $\delta$ | 2   | $\beta$  | b   |

$r$

| $B$ | $D$ | $E$        |
|-----|-----|------------|
| 1   | a   | $\alpha$   |
| 3   | a   | $\beta$    |
| 1   | a   | $\gamma$   |
| 2   | b   | $\delta$   |
| 3   | b   | $\epsilon$ |

$s$

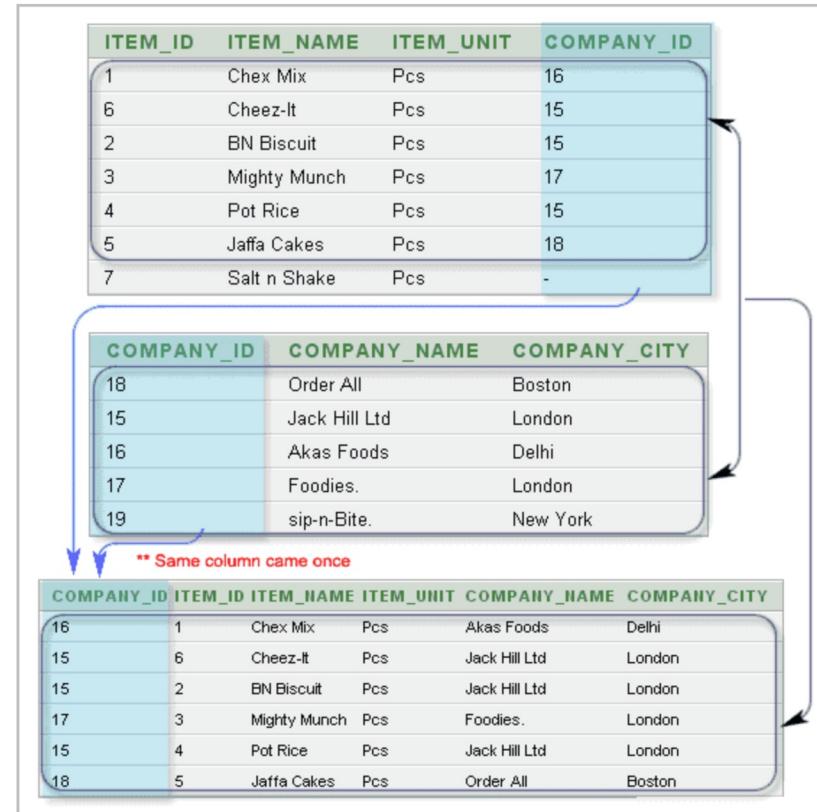
- $r \bowtie s$

| $A$      | $B$ | $C$      | $D$ | $E$      |
|----------|-----|----------|-----|----------|
| $\alpha$ | 1   | $\alpha$ | a   | $\alpha$ |
| $\alpha$ | 1   | $\alpha$ | a   | $\gamma$ |
| $\alpha$ | 1   | $\gamma$ | a   | $\alpha$ |
| $\alpha$ | 1   | $\gamma$ | a   | $\gamma$ |
| $\delta$ | 2   | $\beta$  | b   | $\delta$ |



# Natural Join

- When tables have columns with identical column names.
- Having the same data type.
- Natural join can be used in this case. **No duplicate columns.**
- Example:



```
SELECT LastName , FirstName, OrderID, OrderDate
FROM CUSTOMERS NATURAL JOIN ORDERS
```

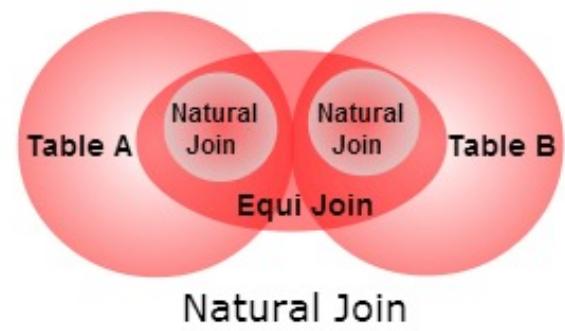
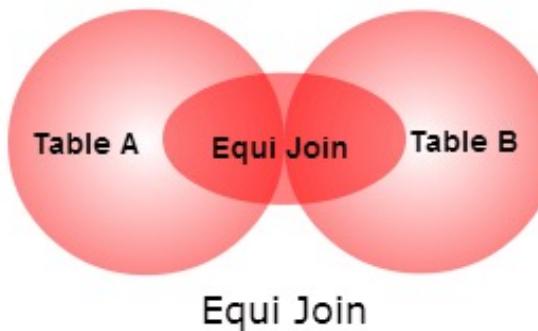
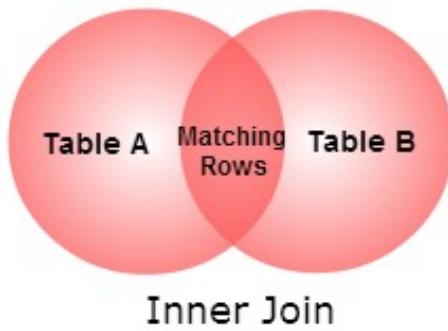


SAINT LOUIS  
UNIVERSITY

# Natural Join vs Inner

## Join

- Natural Join joins two tables based on same attribute name and datatypes.
- Inner Join joins two table on the basis of the column which is explicitly specified in the ON clause.





# JOIN ...USING

- When the tables have some columns with the same column name.
  - Another way is to use:

```
SELECT LastName , FirstName, OrderID, OrderDate  
FROM ORDERS JOIN Customers USING (CustomerID)
```



# JOIN ON

- When the column names are different in the tables.
  - Suppose that the customer id in the customers table is named ID.
  - JOIN....ON can be used.
  - Specify the column to join on.
  - Then:

```
SELECT LastName , FirstName, OrderID, OrderDate  
FROM ORDERS O JOIN CUSTOMERS C  
ON C.ID = O.CustomerID
```



# Inner Join Wrap Up

Traditional: WHERE clause

```
SELECT LastName , FirstName, OrderID, OrderDate  
FROM ORDERS O, CUSTOMERS C WHERE  
O.CustomerID = C.CustomerID
```

NATURAL JOIN

```
SELECT LastName , FirstName, OrderID, OrderDate  
FROM ORDERS NATURAL JOIN CUSTOMERS
```

JOIN.....USING

```
SELECT LastName , FirstName, OrderID, OrderDate  
FROM ORDERS JOIN CUSTOMERS  
USING (CustomerID)
```

JOIN.....ON

```
SELECT LastName , FirstName, OrderID, OrderDate  
FROM ORDERS O JOIN CUSTOMERS C  
ON C.ID = O.CustomerID
```